

Concurrency

Topic of Software Systems (TCS module 2)

Lecturer: Marieke Huisman



WHAT IS CONCURRENCY?

- Most of your programs so far: **single-threaded**
 - Only one execution of the code
 - Also: only one execution to think about
- Now: **multi-threaded** programs (also called **concurrent** or **parallel**)
 - As if multiple programs are running in the same “Java world”
 - Also: more difficult to reason about code correctness
 - Intentional *and* unintentional interaction between threads
- Actually: JVM already has extra threads, for example for garbage collection!

WHAT IS CONCURRENCY?

- Why? More efficient!
 - Computers have multiple processors (and processors have multiple cores)
 - Great for [divide-and-conquer](#) algorithms like mergesort!
 - Threads can be blocked
 - Waiting for disk or network I/O
 - Graphical user interfaces
 - Background operations should not block computations
 - Garbage collection
- Why? More natural!
- How do threads communicate?
 - Via [shared memory](#) and/or [message passing](#) and/or [monitors](#)

WHAT IS CONCURRENCY?

- Thinking about your multi-threaded program




```
// A little experiment with a naive implementation of mergesort  
// using the advanced Java work stealing thread pool to  
// sort an ArrayList of random Integes
```

```
Available hardware threads: 6  
to sort 10000000 integers with Java's default TimSort: 5.455 sec.  
to sort 10000000 integers with parallel mergesort 1: 10.928 sec.  
to sort 10000000 integers with parallel mergesort 2: 6.585 sec.  
to sort 10000000 integers with parallel mergesort 4: 4.646 sec.  
to sort 10000000 integers with parallel mergesort 8: 4.315 sec.
```

- Multiple “workers” (threads) collaborate on a result
- Mergesort: let other threads sort the smaller array concurrently, then merge sequentially
- User interface: some threads handle input, other threads handle output

CHALLENGES WITH CONCURRENCY

- Thinking about your multi-threaded program is **hard**
 - Especially if you really want your program to be fast and efficient
 - Concurrency bugs are **hard to debug**: Heissenbugs
- Two examples of concurrency-specific errors:
 - Accessing objects from multiple threads: **data race** / **race condition**
 - Multiple threads waiting for each other: **deadlock**
- Concurrency errors can be deadly
 - “Therac-25” race condition
 - North American Blackout of 2003



CONCURRENCY IN JAVA

More about concurrency:
Programming Paradigms
(Module 2.4)

- **Thread handling**
 - Thread creation: Implementing Runnable interface
 - Terminated threads can be joined
- Threads **share data**
- Access to data should often be **synchronized** to avoid data races
 - Every object is a lock
 - Synchronized code block
 - Synchronized methods
 - Lock interface
- **Inter-thread communication** about object state
 - Wait-notify
 - More fine-grained: use conditions

Extensive library for
concurrency:
`java.util.concurrent`